

# Report

Gaurav Bhole - 2021113015

## 8.1: Theory:

1. Concept of Soft Prompts: Soft prompts address limitations of discrete text prompts through learnable continuous embeddings that exist in the model's latent space rather than being constrained to actual tokens. Key advantages include:
  - Flexibility in Representation: Soft prompts can learn optimal task-specific representations that may not correspond to any actual tokens in the vocabulary, allowing them to capture more nuanced patterns and relationships than discrete text prompts.
  - Continuous Optimization: Unlike discrete prompts that are limited to existing tokens, soft prompts can be continuously optimized via gradient descent, enabling fine-grained adaptation to the task.
  - Expressive Power: They can encode complex task instructions and conditioning that might be difficult or impossible to express through natural language tokens.
  - Parameter Efficiency: Soft prompts add very few trainable parameters (typically just the prompt embeddings) while allowing effective task-specific adaptation.
  - Task Compositionality: Multiple soft prompts can be combined or interpolated to handle multiple tasks or create new task variations.
2. Scaling and Efficiency in Prompt Tuning: The relationship between model scale and prompt tuning efficiency has important implications:
  - Performance Scaling: As language models get larger, the effectiveness of prompt tuning improves disproportionately. The performance gap between prompt tuning and full fine-tuning narrows significantly with scale.
  - Parameter Efficiency: The number of trainable parameters in prompt tuning remains constant regardless of model size, making it increasingly efficient for larger models. This contrasts with full fine-tuning where parameters grow linearly with model size.
  - Memory Benefits: Prompt tuning requires storing only the prompt parameters, while full fine-tuning needs separate copies of model weights for each task. This becomes crucial for very large models.
  - Future Implications: As models continue to grow in size, prompt tuning may become the default adaptation method due to its combination of strong performance and resource efficiency.

3. Understanding LoRA: Key principles of Low-Rank Adaptation include:
  - Matrix Factorization: LoRA decomposes weight updates into products of lower-rank matrices ( $W + BA$  where  $B$  and  $A$  are low-rank), significantly reducing the number of trainable parameters.
  - Rank-Based Compression: The rank hyperparameter controls the capacity/compression tradeoff, allowing flexible scaling of adaptation complexity.
  - Parameter Freezing: Original model weights remain frozen while only the low-rank adaptation matrices are trained.
  - Training Efficiency: By training fewer parameters, LoRA reduces memory requirements and training time while maintaining performance comparable to full fine-tuning.
  
4. Theoretical Implications of LoRA: The success of LoRA has several important theoretical implications:
  - Low-Rank Structure: The effectiveness of low-rank updates suggests that many model adaptations naturally lie in a low-dimensional subspace, challenging the need for full parameter fine-tuning.
  - Implicit Regularization: The low-rank constraint acts as a form of regularization, potentially improving generalization by restricting the space of possible adaptations.
  - Model Capacity: LoRA demonstrates that model capacity can be effectively expanded through targeted, low-dimensional updates rather than full parameter tuning.
  - Optimization Landscape: The success of low-rank updates suggests that the optimization landscape for model adaptation may be more structured than previously thought.
  - Modularity: The ability to achieve strong performance with modular, low-rank updates suggests possibilities for more efficient and composable model adaptation strategies.

## 8.3: Analysis Report: Comparison of Fine-tuning Methods for GPT-2 on Summarization Task

### Training Configuration

Common hyperparameters across all methods:

- Base Model: GPT-2 Small
- Number of Epochs: 20
- Batch Size: 8
- Initial Learning Rate: 1e-4
- Optimizer: AdamW
- Loss Function: CrossEntropyLoss
- Learning Rate Schedule: Linear with warmup (100 steps)
- Dataset: 10% of CNN/Daily Mail dataset
- Maximum sequence lengths: 512 (input), 128 (summary)

Method-specific configurations:

1. Prompt Tuning:
  - Soft prompt length: 20 tokens
  - Only prompt embeddings trainable
  - Random initialization with range  $\pm 0.5$
2. LoRA:
  - Rank: 8
  - Alpha: 32
  - Applied to attention layers
  - Initialization: Zero for B matrices, Normal for A matrices
3. Traditional Fine-tuning:
  - Only last transformer block and LM head trainable
  - All other parameters frozen

### Performance Analysis

#### 1. Training and Validation Loss

All three methods showed similar loss patterns:

- Initial rapid decrease in training loss
- Stabilization after ~5 epochs
- Slight divergence between training and validation loss

Best Validation Losses:

1. Prompt Tuning: 0.9863 (epoch 5)
2. Traditional Fine-tuning: 0.9998 (epoch 4)
3. LoRA: 1.0006 (epoch 7)

Prompt Tuning achieved the lowest validation loss, suggesting better generalization despite having fewer trainable parameters.

## 2. ROUGE Scores

All methods achieved similar ROUGE scores:

ROUGE-1: ~0.056 (5.6%)

ROUGE-2: ~0.031 (3.1%)

ROUGE-L: ~0.049 (4.9%)

The similarity in ROUGE scores across methods suggests that all approaches were equally effective at learning the summarization task, despite their different mechanisms.

## 3. Resource Usage

GPU Memory Usage:

- LoRA: 94.24%
- Traditional Fine-tuning: 94.24%
- Prompt Tuning: 89.43%

CPU Usage (relatively consistent across methods):

- LoRA: 7.62%
- Traditional Fine-tuning: 7.74%
- Prompt Tuning: 7.52%

Prompt Tuning showed slightly lower GPU memory usage while maintaining comparable performance.

## 4. Learning Dynamics

Training Loss Patterns:

- Prompt Tuning: Most aggressive initial decrease (1.29 → 0.83)
- LoRA: Moderate decrease (1.05 → 0.95)
- Traditional Fine-tuning: Similar to LoRA (1.06 → 0.95)

All methods showed stable learning with minimal oscillation in validation loss.

## 5. Parameter Efficiency

Total Parameters vs. Trainable Parameters:

- Base Model: ~124M parameters
- Prompt Tuning: Added minimal parameters (soft prompts only)
- LoRA: Added low-rank matrices
- Traditional Fine-tuning: No additional parameters

## Key Findings

1. **Effectiveness:** All three methods achieved comparable performance metrics, suggesting they are all viable approaches for the summarization task.
2. **Efficiency:**
  - Prompt Tuning showed the best parameter efficiency while achieving slightly better validation loss
  - LoRA provided a good balance between performance and parameter count
  - Traditional Fine-tuning was competitive despite only updating the last layers
3. **Stability:** All methods demonstrated stable training with minimal overfitting, as evidenced by the relatively small gap between training and validation losses.
4. **Resource Usage:** Prompt Tuning had a slight edge in GPU memory efficiency, while CPU usage was comparable across methods.

## Recommendations

1. For resource-constrained environments: Choose Prompt Tuning due to its lower memory footprint and competitive performance.
2. For best balance of performance and complexity: Use LoRA, as it provides comparable results with a clean implementation approach.
3. For simplicity and reliability: Traditional Fine-tuning remains a solid choice, showing robust performance with minimal complexity.

## Limitations

1. Relatively low ROUGE scores across all methods suggest room for improvement in the summarization quality.
2. The training was conducted on only 10% of the dataset, which might not represent full potential performance.